

AI-Based Business Information Systems

Generative AI



Prof. Dr. Ulrich Gnewuch

Lecture

AI-Enabled Business Capabilities

AI-Enabled Innovation

AI-Enabled Insights & Decisions

AI-Enabled Engagement

AI-Enabled Automation

AI Technologies & Trends

AI Ethics & Ethical AI

Generative AI

Explainable AI

Conversational AI

Foundations

Introduction to AI in Business
& Information Systems

Design & Management of AI-
Based Information Systems

Exercise


Exercise 4:
Generative AI &
Innovation

Exercise 3:
Explainable AI
Techniques

Exercise 2:
Human-Centered
Chatbot Design

Exercise 1:
Robotic Process
Automation Case Study

Industry Talk
ZF Group



Mentimeter




RECAP FROM LAST LECTURE:

- Which innovation type best describes the replacement of human workers with assembly robots in a car factory?
- Which statements about 'creativity' are true?
- What are key challenges associated with the use of AI in innovation?



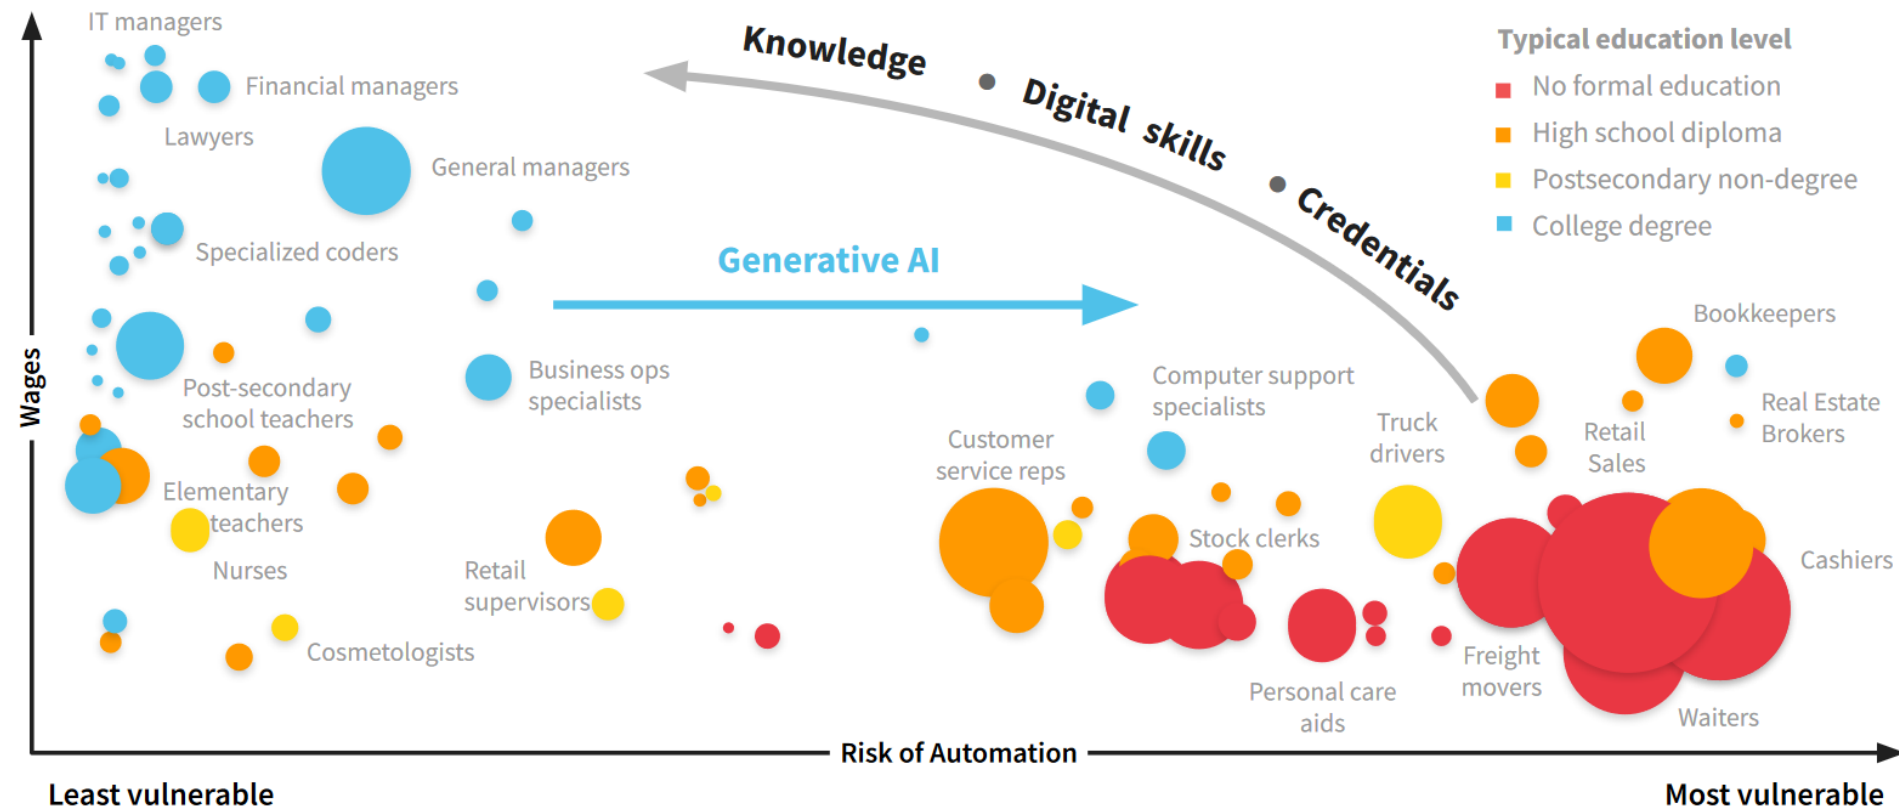
- Explain what generative AI is and how it differs from “traditional” AI
- Describe how foundation models can be adapted to specific tasks and domains
- Explain key prompting strategies and apply them effectively
- Discuss the challenges and risks of generative AI



Gartner 2024

Low-skilled jobs are at risk of automation

AI has the potential to impact an entire new class of knowledge workers, unleashing a new wave of reskilling imperatives




https://s27.q4cdn.com/928340662/files/doc_financials/2023/q1/COUR_Presentation_Q1-2023.pdf



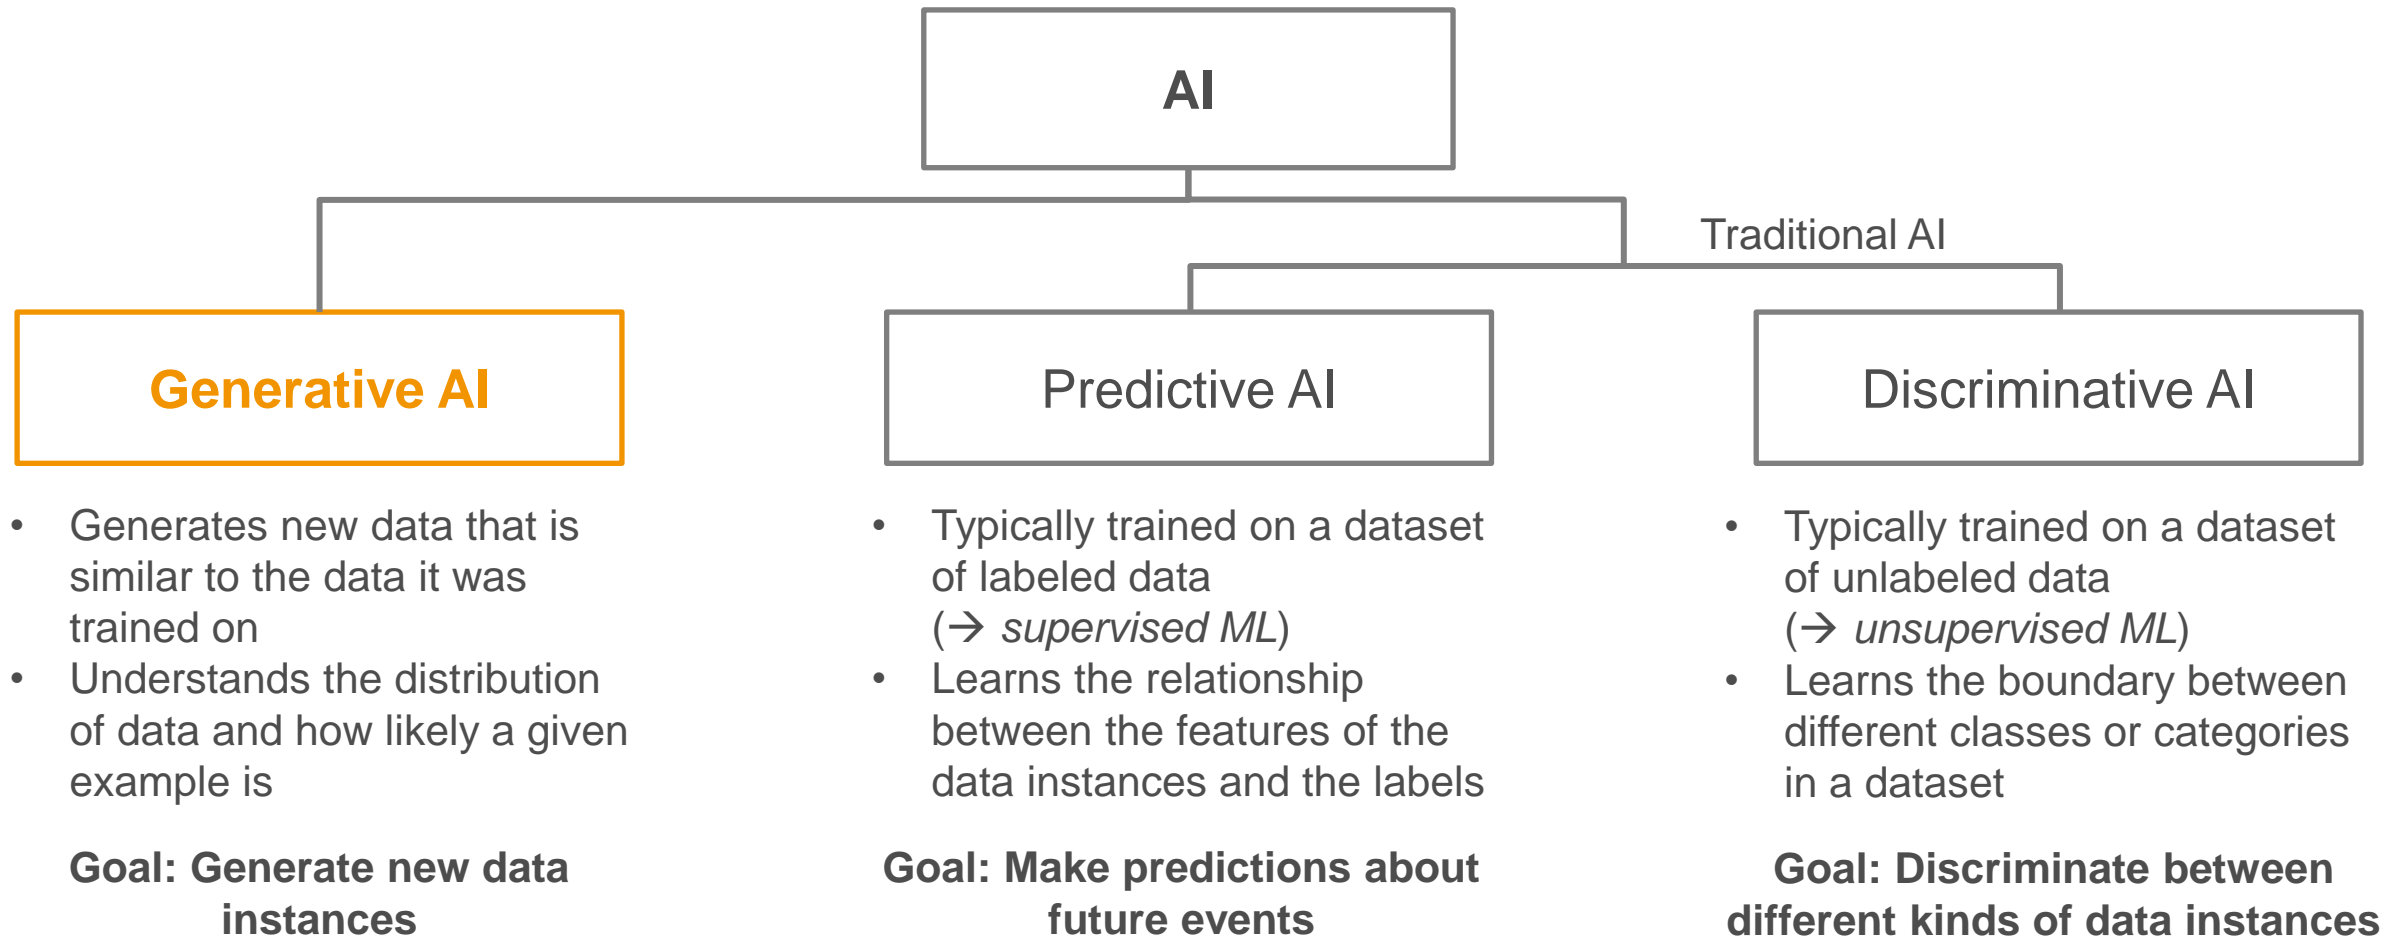
Generative AI is a type of AI that generates seemingly new, meaningful content, such as text, images, audio, or video, from training data.

- Generative AI relies on statistical models created from existing content
- When given a prompt, generative AI uses this statistical model to generate new content



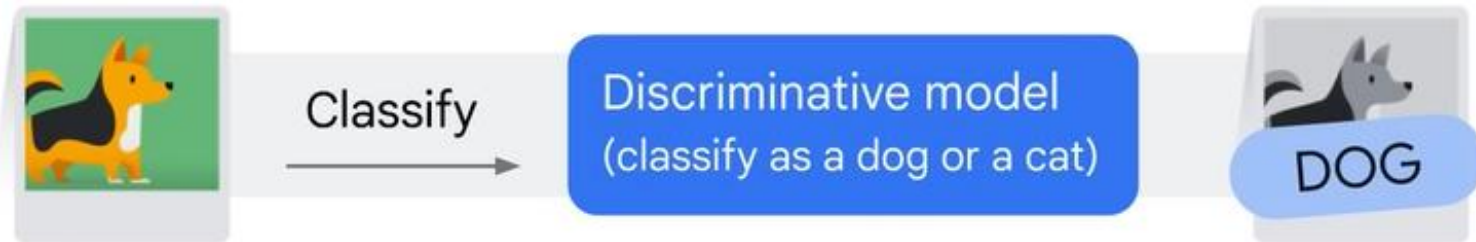
Feuerriegel et al. 2024

Generative vs. “Traditional” AI




Generative vs. Traditional AI: Example

**Traditional AI
(Discriminative)**

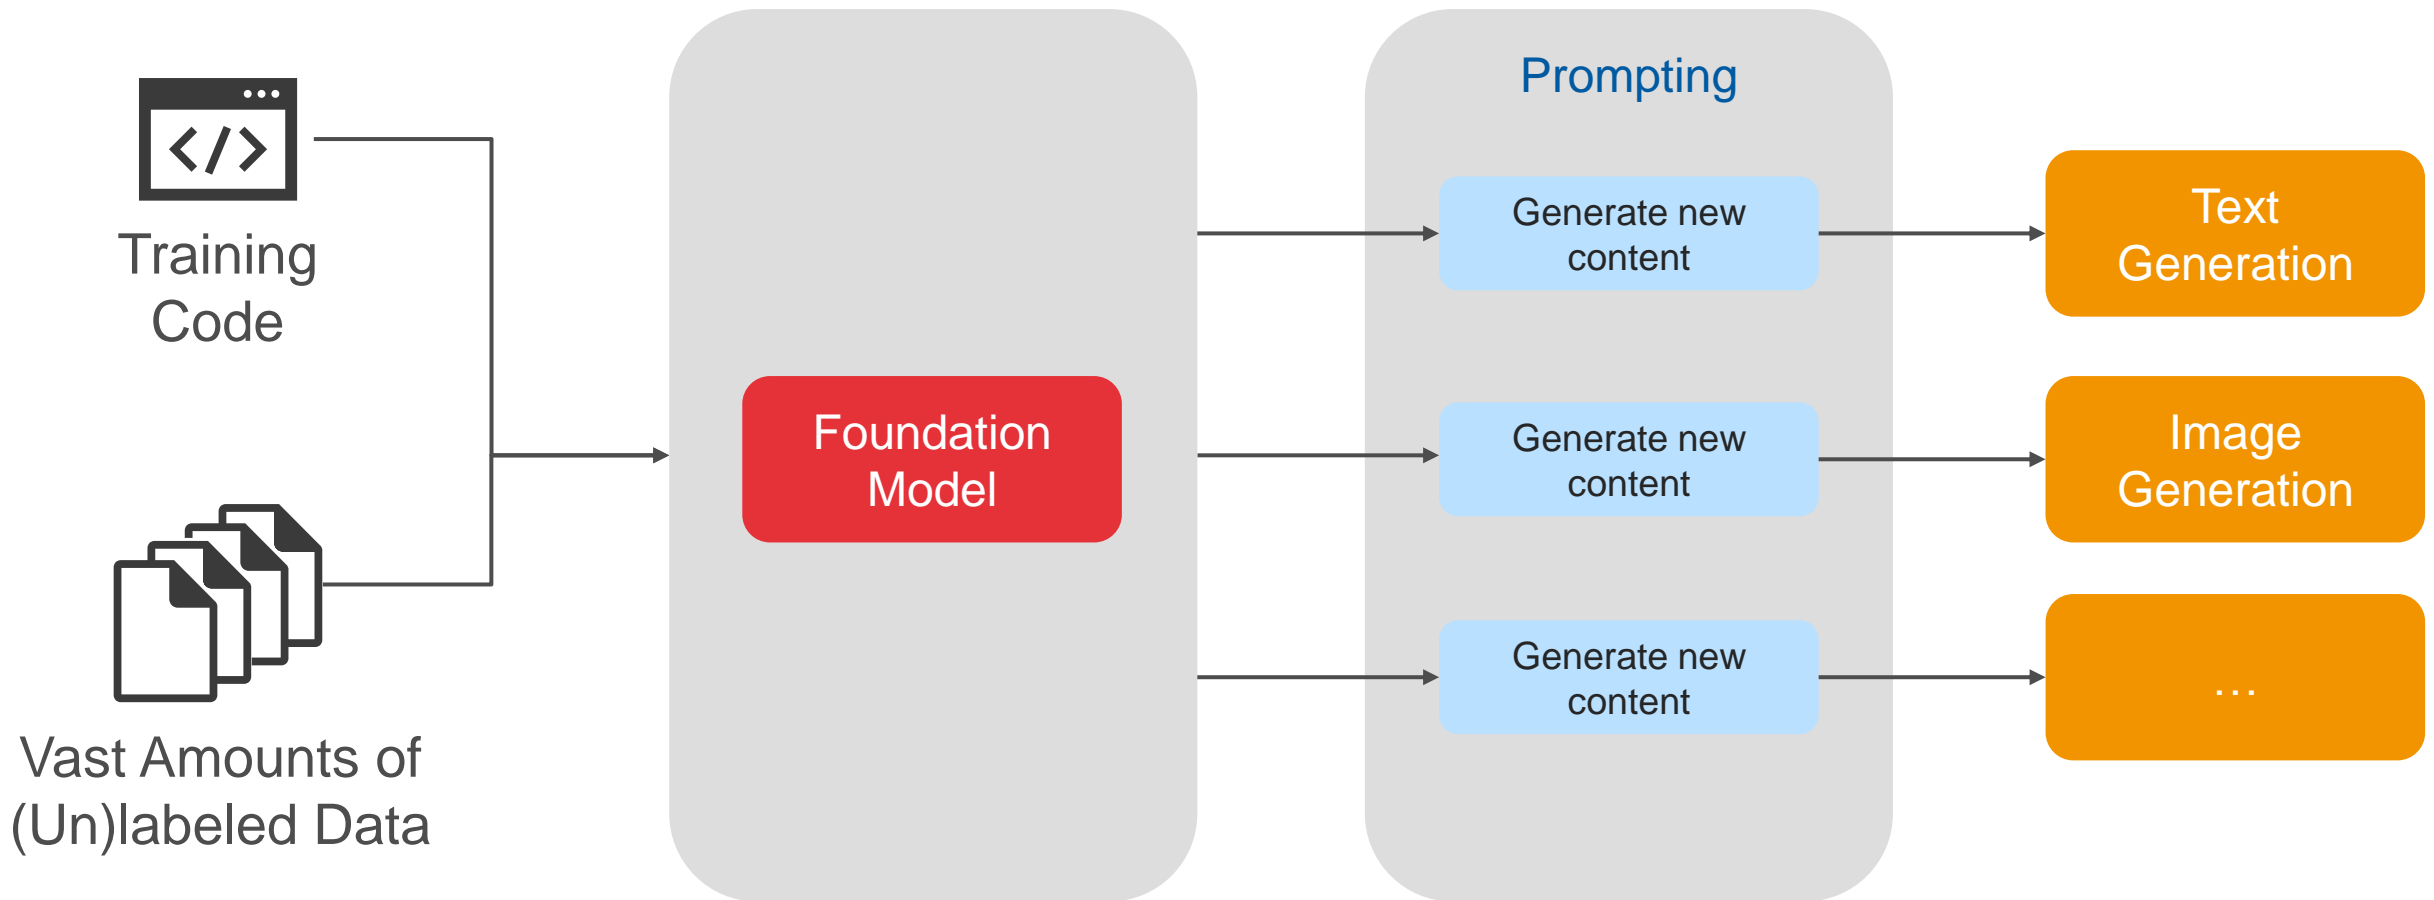


Generative AI





Based on: Google 2023




Based on: Google 2023

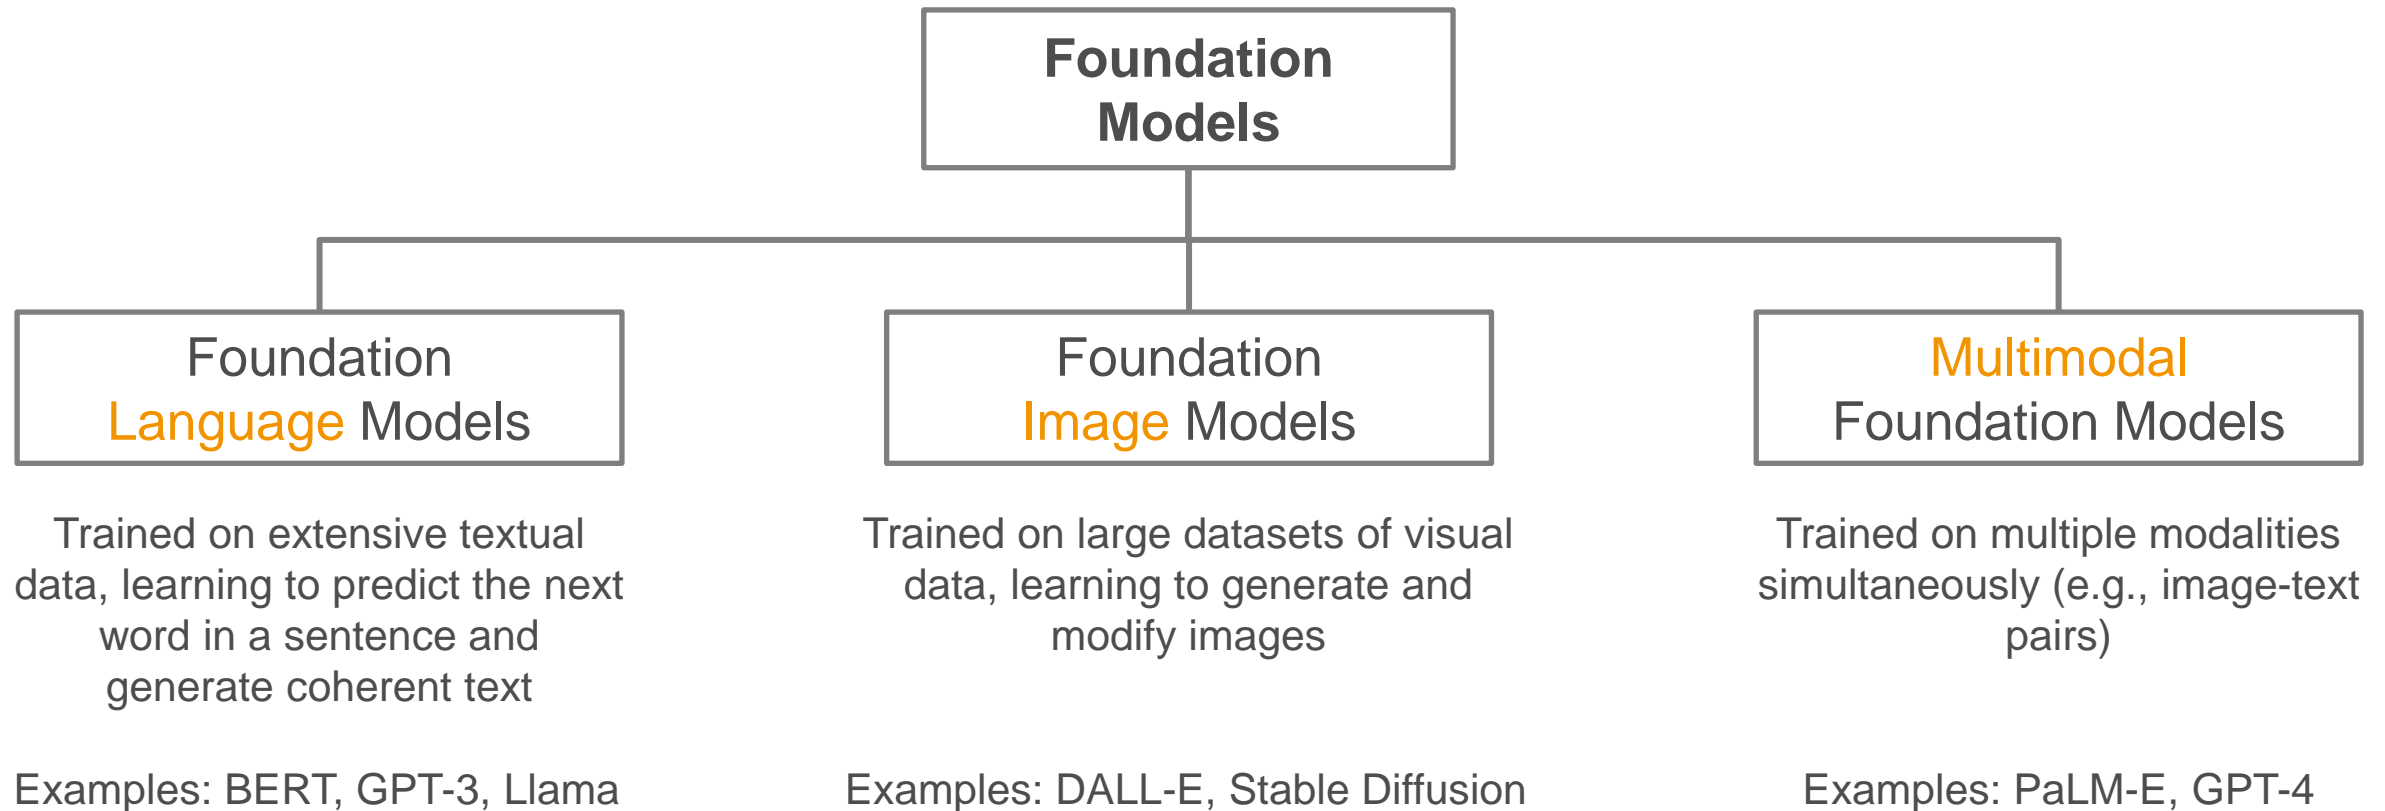


A foundation model is an AI model trained with a large amount of data using self-supervision at scale, displays significant generality, and is capable of performing a wide range of distinct tasks.

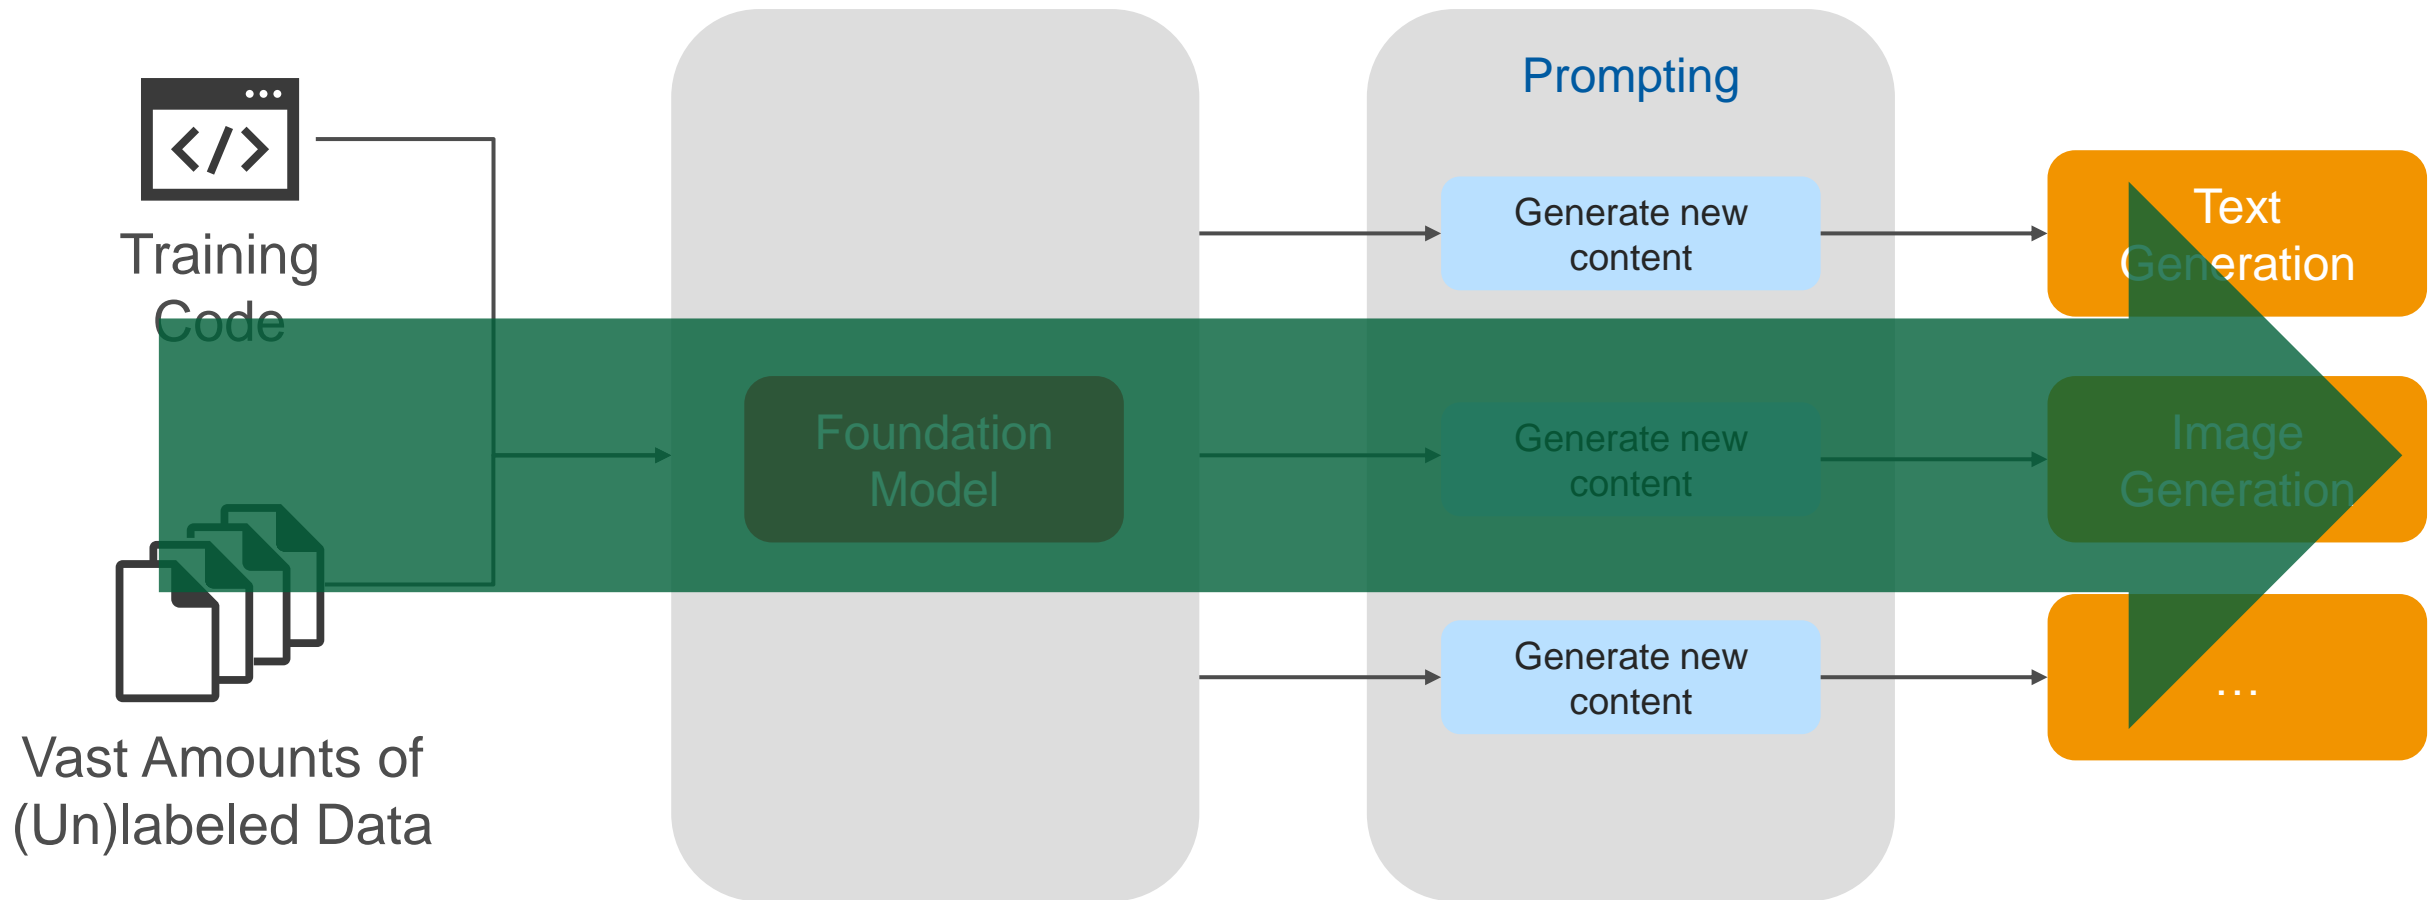
- These models are sometimes referred to as 'general-purpose AI (GPAI) models' or 'large X models' (LxM)
- Example: GPT-4 is the fourth generation of OpenAI's foundation model



EU AI Act (Article 3, 63)

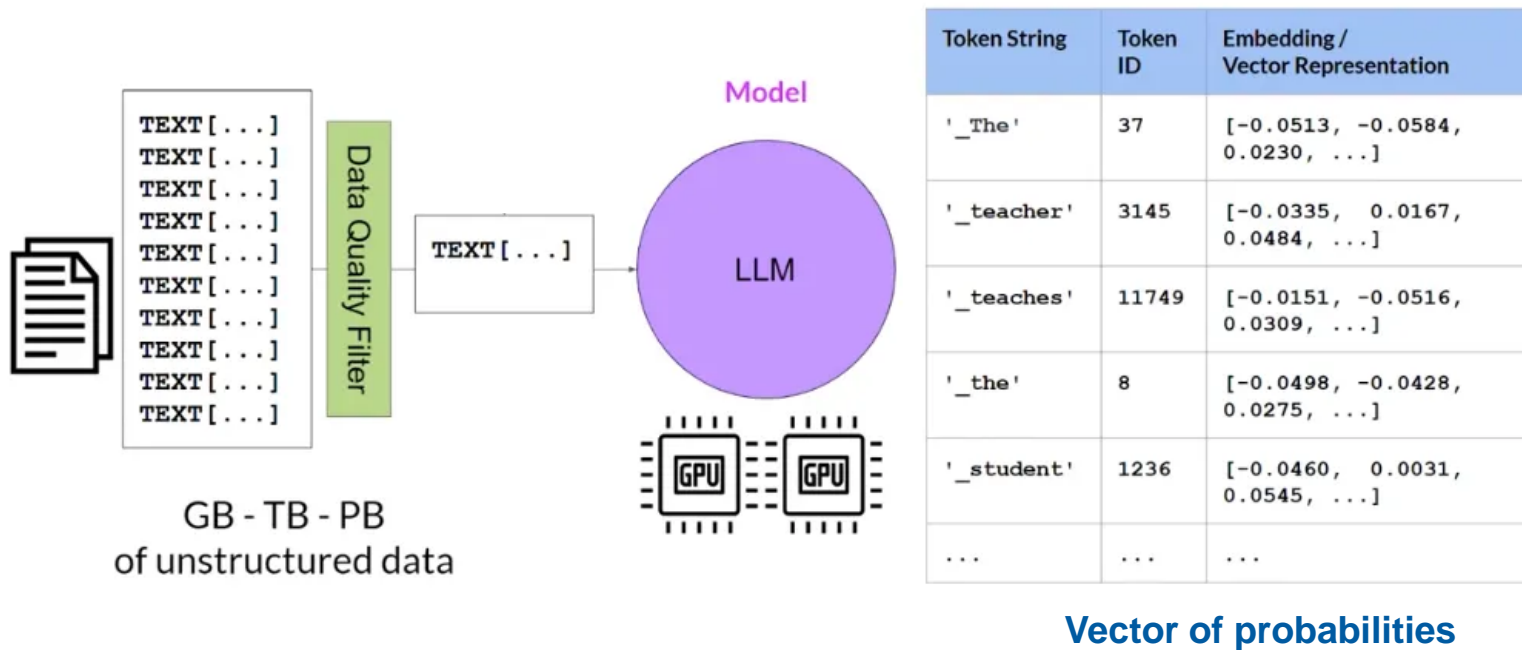


Generative AI: How does it work?



Based on: Google 2023

- From a very large training dataset, a model learns the **statistical distributions of tokens**
- Tokens can be “chunks” of words, punctuation marks, pixels, etc.



Faklaris 2023, Stollnitz 2023

- When given a prompt, the model converts the prompt into tokens and then **analyzes what is likely to come next**, based on the tokens in its own dataset


“Here’s a fragment of text.

*Tell me how this fragment might **<continue on in this language, or suggest a particular image>**.*

*According to your model of the statistics of **<human language, or human-handled images>**, what **<words, or pixels>** are likely to come next?”*

- When given a prompt, the model converts the prompt into tokens and then analyzes what is likely to come next, based on the tokens in its own dataset
- It then **generates a tokenized output**

The best thing about AI is its ability



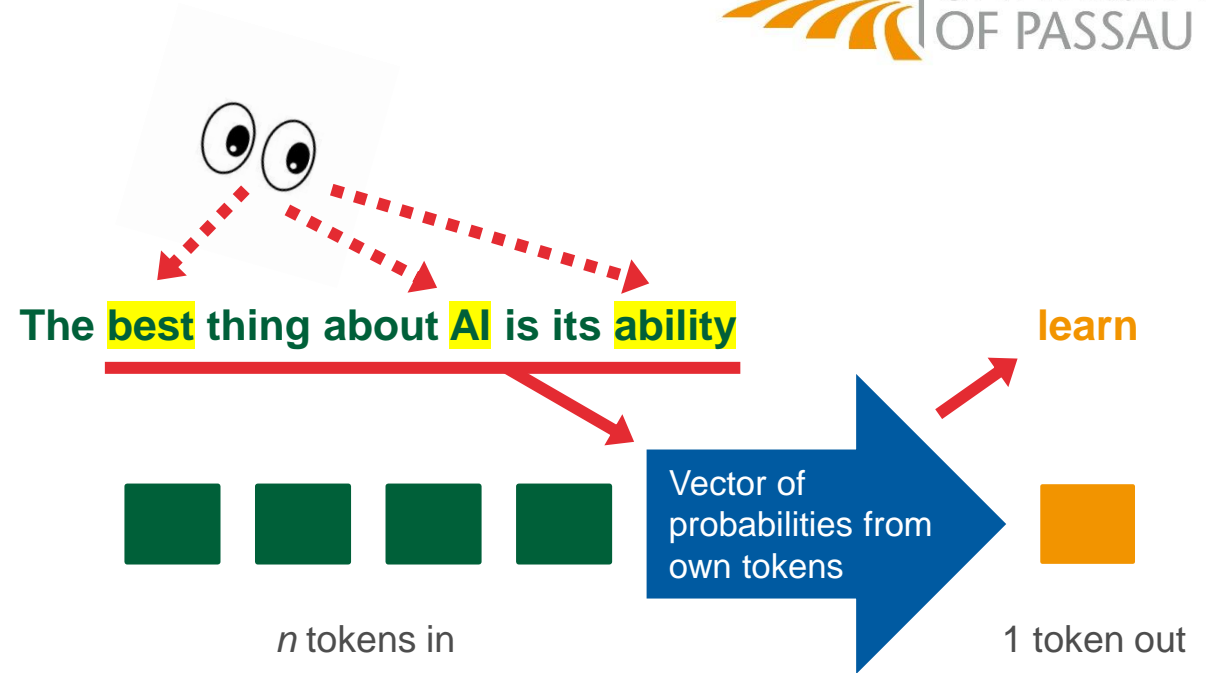
- With each output, the model **keeps reanalyzing** the probabilities to decide the next token

The best thing about AI is its ability to

learn



- Transformers (the “T in “GPT”) know how to direct attention to specific parts of the input to guide their selection of the next token
- The transformer architecture was a major breakthrough in the field of natural language processing (see paper “*Attention is all you need*”; Vaswani et al. 2017)




- The model can provide different tokens to the same inputs:

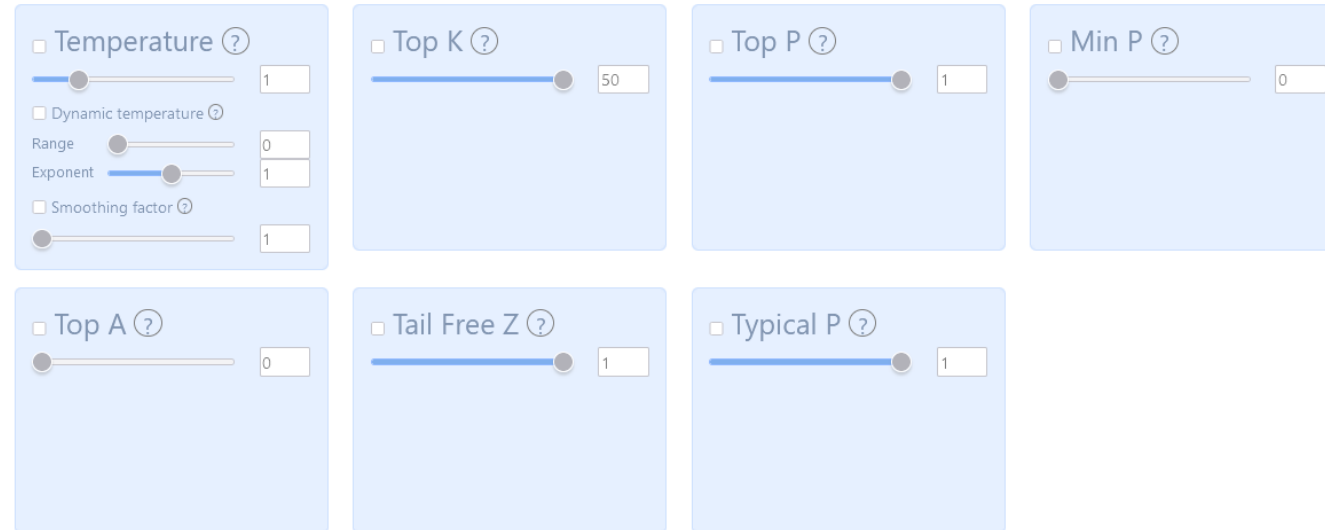
learn	4.5%
predict	3.5%
make	3.2%
understand	3.1%
do	2.9%

The best thing about AI is its ability to

predict



No, Johnny, Steve didn't jump off a bridge. That was just a figure of speech. I'm glad you're ____



The interface contains seven control panels for LLM sampling parameters:

- Temperature**: A slider set to 1, with checkboxes for Dynamic temperature, Range, Exponent, and Smoothing factor.
- Top K**: A slider set to 50.
- Top P**: A slider set to 1.
- Min P**: A slider set to 0.
- Top A**: A slider set to 0.
- Tail Free Z**: A slider set to 1.
- Typical P**: A slider set to 1.

Token	Probability	
concerned	27,837 %	<div></div>
interested	22,259 %	<div></div>
taking	10,004 %	<div></div>
curious	7,804 %	<div></div>

<https://chatbotresearch.org/llm-sampling/index.xhtml>

- LLMs can use different strategies to choose the next token
- This process is known as “sampling”
- Example based on Meta’s LLaMA model
 - Source:
<https://github.com/Artefact2/llm-sampling>



<https://shorturl.at/Midca>

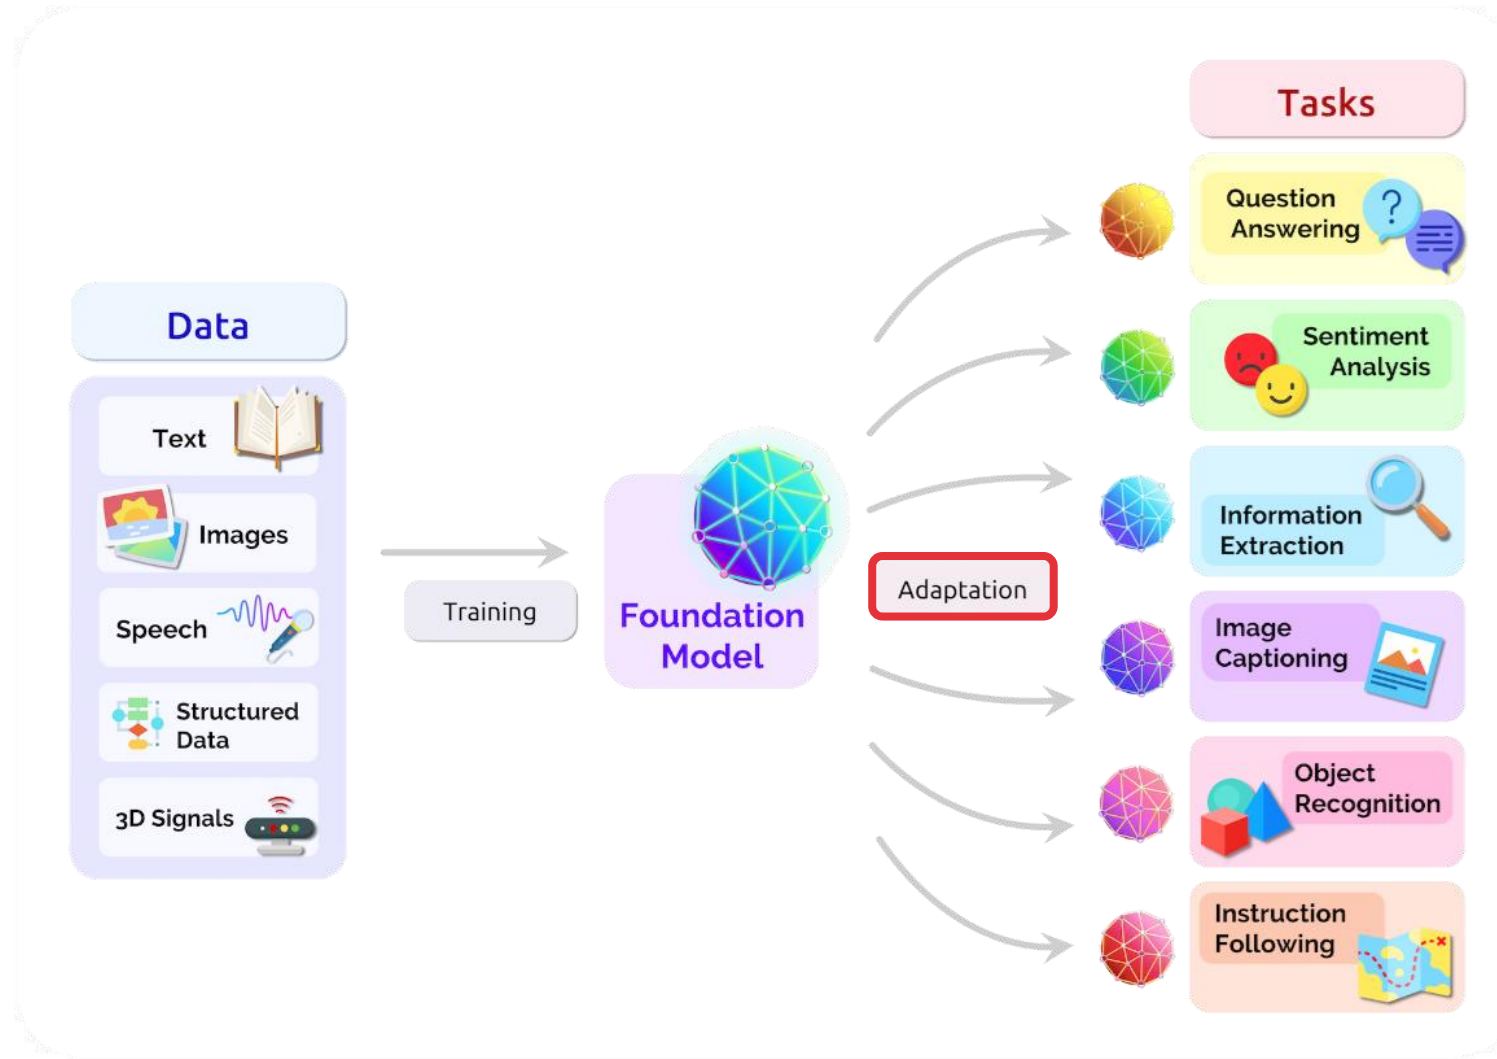
LLM Demo & Token Sampling

Open the interactive LLM demo, explore how tokens are sampled under different parameter settings, and answer the following questions:

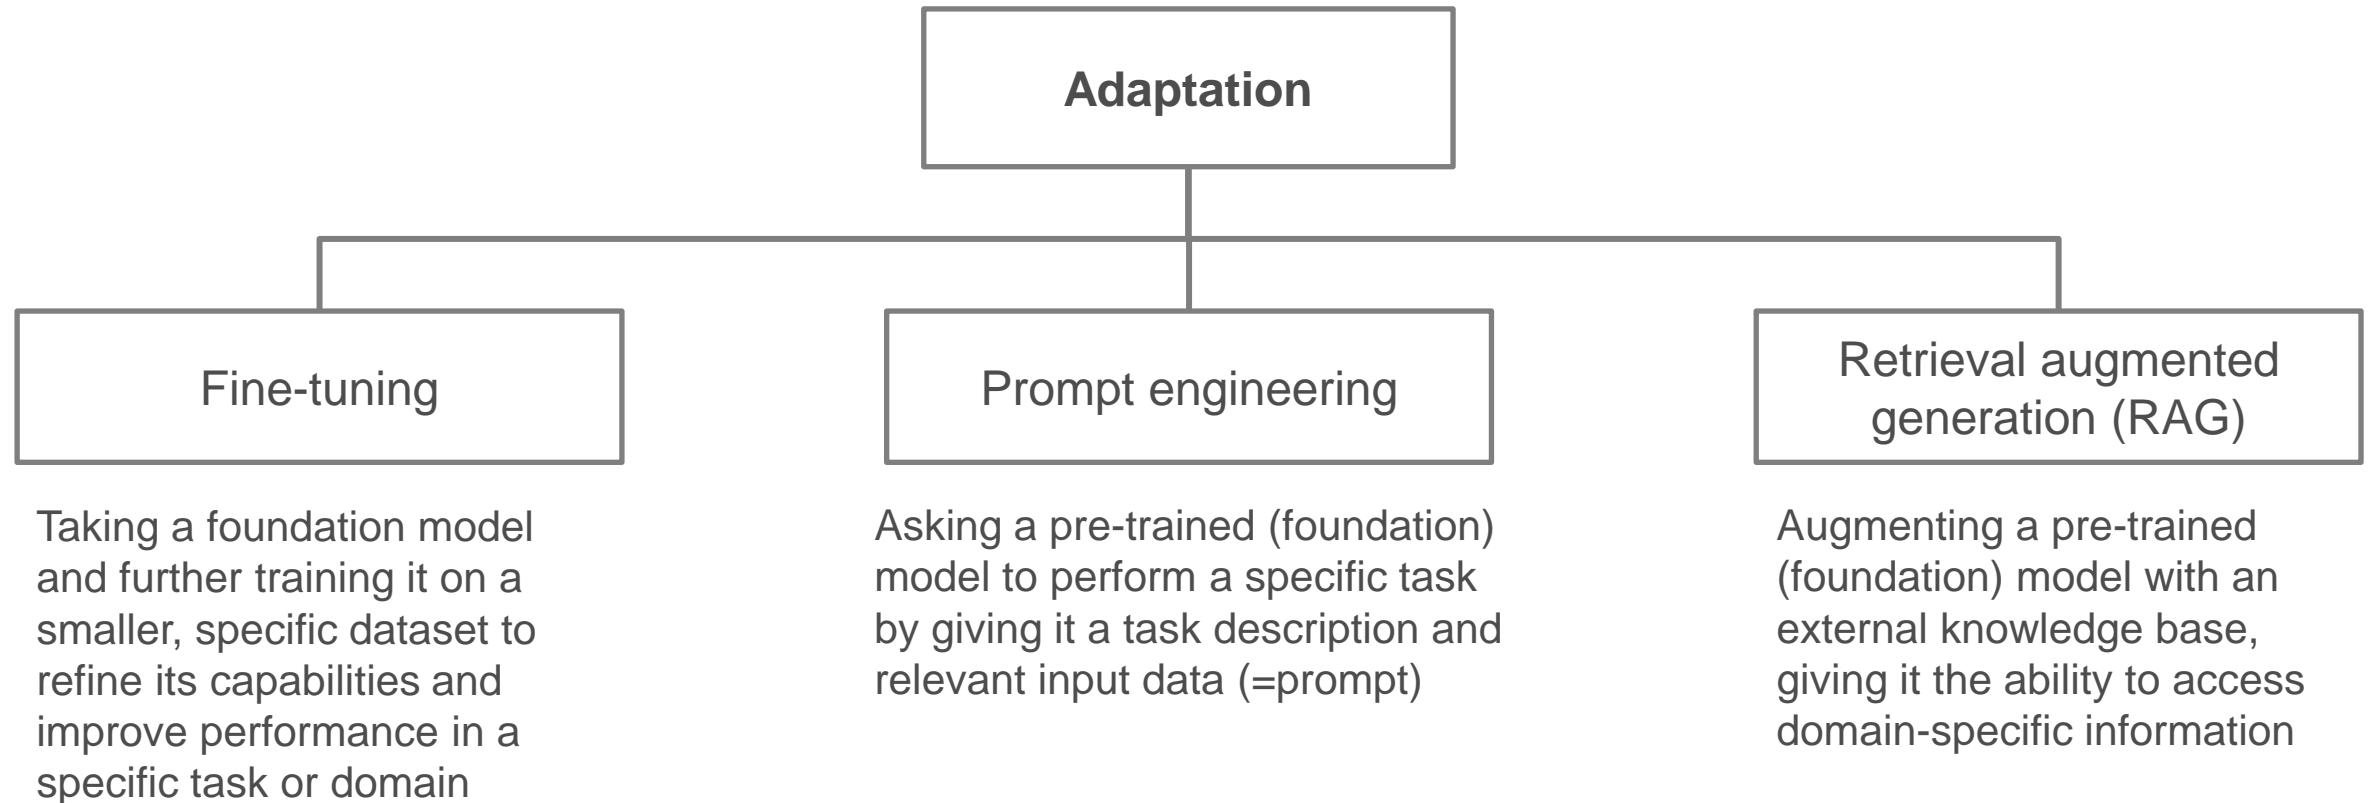
1. What happens when the temperature is zero? What happens when it is set to its maximum value?
2. Why do you think the temperature is typically set to a value between 0 and 1 by OpenAI, Anthropic, and other LLM providers?

→ Discuss these questions with a partner for **~5 minutes** and be ready to share your answers

Adapting Foundation Models to Specific Tasks and Domains



Bommasani et al. 2021



- Fine-tuning is a powerful technique but requires sufficiently large annotated training data which is often costly and hard to obtain
- In practice, combinations of different techniques are often used
- New techniques are constantly being developed (e.g., parameter-efficient fine-tuning; PEFT)



Bommasani et al. 2021

Fine-Tuning: Example

Introducing BloombergGPT, Bloomberg's 50-billion parameter large language model, purpose-built from scratch for finance

NEW YORK - Bloomberg today released a research paper detailing the development of BloombergGPT™, a new large-scale generative artificial intelligence (AI) model. This large language model (LLM) has been specifically trained on a wide range of financial data to support a diverse set of natural language processing (NLP) tasks within the financial industry.

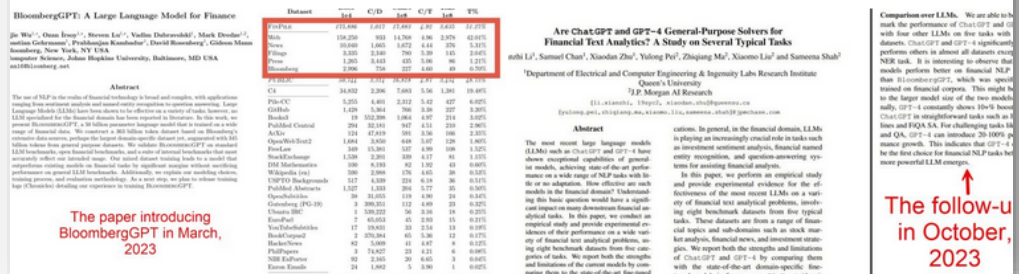
<https://www.bloomberg.com/company/press/bloomberggpt-50-billion-parameter-llm-tuned-finance/>

 **Ethan Mollick**  • 2nd
Associate Professor at The Wharton School. Author of Co-Intelli...
8mo • 

This remains one of the most consequential experiments in AI: Bloomberg over \$10M training a GPT-3.5 class AI on their own financial data last year

...only to find that GPT-4 8k, the AI available to billions of people around the world, and without specialized finance training, beat it on almost all finance tasks!

There was a moment that we thought proprietary data would let organizations train specialized AIs that could compete with frontier models. It turns out that probably isn't going to happen. The largest frontier models are just much better at most complex tasks than smaller models.



BloombergGPT: A Large Language Model for Finance

Dataset	BLEU	C/D	C/P	C/T	F1
CoNLL	0.74	0.75	0.76	0.77	0.78
Wiki	0.75	0.76	0.77	0.78	0.79
News	0.76	0.77	0.78	0.79	0.80
Finance	0.77	0.78	0.79	0.80	0.81
Reddit	0.78	0.79	0.80	0.81	0.82
Stack	0.79	0.80	0.81	0.82	0.83
ArXiv	0.80	0.81	0.82	0.83	0.84
OpenWebText2	0.81	0.82	0.83	0.84	0.85
News	0.82	0.83	0.84	0.85	0.86
Reddit	0.83	0.84	0.85	0.86	0.87
Stack	0.84	0.85	0.86	0.87	0.88
ArXiv	0.85	0.86	0.87	0.88	0.89
OpenWebText2	0.86	0.87	0.88	0.89	0.90
News	0.87	0.88	0.89	0.90	0.91
Reddit	0.88	0.89	0.90	0.91	0.92
Stack	0.89	0.90	0.91	0.92	0.93
ArXiv	0.90	0.91	0.92	0.93	0.94
OpenWebText2	0.91	0.92	0.93	0.94	0.95
News	0.92	0.93	0.94	0.95	0.96
Reddit	0.93	0.94	0.95	0.96	0.97
Stack	0.94	0.95	0.96	0.97	0.98
ArXiv	0.95	0.96	0.97	0.98	0.99
OpenWebText2	0.96	0.97	0.98	0.99	1.00

Are ChatGPT and GPT-4 General-Purpose Solvers for Financial Text Analysis? A Study on Several Typical Tasks

Authors: Li, Samuel Chai*, Xiaodan Zhu*, Yuhang Pu*, Zhaoyang Ma*, Xianxian Liu* and Samanta Shah*

**Department of Electrical and Computer Engineering & Ingersoll Labs Research Institute, Queen's University*

Abstract

The most recent large language models (LLMs) such as ChatGPT and GPT-4 have shown exceptional capabilities of general text analysis, achieving state-of-the-art performance on a wide range of NLP tasks with little or no adaptation. However, none of the models in the financial domain. Understanding the basic question would have a significant impact on many downstream financial analysis tasks. In this paper, we conduct an empirical study and provide experimental evidence of their performance on a wide variety of financial text analytical problems, involving eight benchmark datasets from five typical tasks. These datasets are from a range of financial topics and sub-domains such as stock market analysis, financial news, and investment strategies. We report both the strengths and limitations of ChatGPT and GPT-4 by comparing them with the state-of-the-art domain-specific financial models. Our results show that ChatGPT and GPT-4 are not yet ready to replace the state-of-the-art financial models.

The paper introducing BloombergGPT in March, 2023

The follow-up in October, 2023

https://www.linkedin.com/posts/emollick_this-remains-one-of-the-most-consequential-activity-7176398465004896256-Qjx-/



Prompting is the process of using carefully crafted phrases or templates to help a pre-trained (foundation) model accomplish a specific, downstream task.

- Targeted prompting can optimize the performance and relevance of model outputs to produce high-quality results
- The quality of the prompt determines the quality of the output!

**GARBAGE IN,
GARBAGE OUT**



Brown et al. 2020

Prompting Technique	Description	Example
Zero-shot prompting	Relying solely on a model's pre-trained information to answer a given prompt	<ul style="list-style-type: none">• <i>Write a summary of the following text: [...]</i>• <i>Generate ideas for a birthday present.</i>
Few-shot prompting	Providing examples in the prompt, giving the model more context to improve its performance	<i>Country: Germany</i> <i>Capital city: Berlin</i> <i>Country: Argentina</i> <i>Capital city:</i>
Chain-of-thought prompting	Nudging a model to produce intermediate reasoning steps, thus improving the results on complex reasoning tasks	Two ways: <ul style="list-style-type: none">• Use few-shot prompting by illustrating examples with detailed answers to questions• Add instructions such as “<i>Let's think step by step</i>”

<https://huggingface.co/docs/transformers/main/tasks/prompting#advanced-prompting-techniques>

- Prompting can seem easy, but designing effective prompts is challenging, especially for non-experts
- Common challenges:
 - Expectations stemming from human-to-human communication (e.g., politeness)
 - Tendency to overgeneralize
 - Biases towards giving instruction over depicting examples

Why Johnny Can't Prompt: How Non-AI Experts Try (and Fail) to Design LLM Prompts

J.D. Zamfirescu-Pereira
zamfi@berkeley.edu
UC Berkeley
Berkeley, CA, USA

Bjoern Hartmann
bjoern@eecs.berkeley.edu
UC Berkeley
Berkeley, CA, USA

Richmond Wong
rwong34@gatech.edu
Georgia Institute of Technology
Atlanta, Georgia, USA

Qian Yang
qianyang@cornell.edu
Cornell University
Ithaca, NY, USA

ABSTRACT

Pre-trained large language models ("LLMs") like GPT-3 can engage in fluent, multi-turn instruction-taking out-of-the-box, making them attractive materials for designing natural language interactions. Using natural language to steer LLM outputs ("prompting") has emerged as an important design technique potentially accessible to non-AI experts. Crafting effective prompts can be challenging, however, and prompt-based interactions are brittle. Here, we explore whether non-AI experts can successfully engage in "end-user prompt engineering" using a design probe—a prototype LLM-based chatbot design tool supporting development and systematic evaluation of prompting strategies. Ultimately, our probe participants explored prompt designs opportunistically, not systematically, and struggled in ways echoing end-user programming systems and interactive machine learning systems. Expectations stemming from human-to-human instructional experiences, and a tendency to over-generalize, were barriers to effective prompt design. These findings have implications for non-AI-expert-facing LLM-based tool design and for improving LLM-and-prompt literacy among programmers and the public, and present opportunities for further research.

CCS CONCEPTS

• Human-centered computing → Empirical studies in interaction design; • Computing methodologies → Natural language processing.

KEYWORDS

language models, end-users, design tools

ACM Reference Format:

J.D. Zamfirescu-Pereira, Richmond Wong, Bjoern Hartmann, and Qian Yang. 2023. Why Johnny Can't Prompt: How Non-AI Experts Try (and Fail) to Design LLM Prompts. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems (CHI '23)*, April 29–28, 2023, Hamburg, Germany. ACM, New York, NY, USA, 21 pages. <https://doi.org/10.1145/3545458.3581388>



This work is licensed under a Creative Commons Attribution-NonCommercial International 4.0 License.

CHI '23, April 29–28, 2023, Hamburg, Germany
© 2023 Copyright held by the owner(s).
ACM ISBN 978-1-4503-9421-5/23/04.
<https://doi.org/10.1145/3545458.3581388>

1 INTRODUCTION

The idea of instructing computers in natural language has fascinated researchers for decades, as it promises to make the power of computing more customizable and accessible to people without programming training [4]. The combination of pre-trained large language models (LLMs) and prompts brought renewed excitement to this vision. Recent pre-trained LLMs (e.g., GPT-3 [8], ChatGPT [1]) can engage in fluent, multi-turn conversations out-of-the-box, substantially lowering the data and programming-skill barriers to creating passable conversational user experiences [7]. People can improve LLM outputs by prepending prompts—textual instructions and examples of their desired interactions—to LLM inputs. Prompts directly bias the model towards generating the desired outputs, raising the ceiling of what conversational UX is achievable for non-AI experts. In the past two years, social media platforms have witnessed an explosion of posts showing the results of lay peoples' experimentation with LLMs for question answering, creative dialogue writing, writing code, and more. This excitement around LLMs and prompting is propelling a rapidly growing set of LLM-powered applications [23] and prompt design tools [3, 20, 32].

Yet despite widespread excitement, surprisingly little is known about how non-experts intuitively approach designing prompts with LLM-and-prompt-based tools, and how effective they are in doing so. While prompting LLMs can appear effortless, designing effective prompting strategies requires identifying the contexts in which these LLMs' errors arise, devising prompting strategies to overcome them, and systematically assessing those strategies' effectiveness. These tasks fall on so-called "prompt engineers"—the designers, domain experts, and any other end-user or professional attempting to improve an LLM's output—and are challenging tasks even for LLM experts, as well as topics of ongoing research in Natural Language Processing (NLP) [7, 30, 42]. Prompt design tools to date have focused on supporting professional programmers [45] and NLP practitioners [42], rather than non-AI experts, non-programmers, and other potential end-users of these systems.

In this work, we investigate how non-AI experts intuitively approach prompt design when designing LLM-based chatbots, with an eye towards how non-AI-expert-facing design tools might help. Specifically, we investigate these questions in the context of designing an instructional chatbot, that is, a chatbot that walks the user through an activity (e.g., cooking a recipe, fixing a wifi connection) while answering user questions

Company	Guideline	Link
OpenAI	OpenAI's Prompt engineering strategies	https://platform.openai.com/docs/guides/prompt-engineering
Hugging Face	Best practices of LLM prompting	https://huggingface.co/docs/transformers/main/tasks/prompting#best-practices-of-llm-prompting
Google	Tips to enhance your prompt-engineering abilities	https://cloud.google.com/blog/products/application-development/five-best-practices-for-prompt-engineering
Midjourney	Prompting Notes	https://docs.midjourney.com/docs/prompts


- Start with a simple and short prompt, and iterate from there.
- Put the instructions at the beginning of the prompt, or at the very end. When working with large context, models apply various optimizations to prevent Attention complexity from scaling quadratically. This may make a model more attentive to the beginning or end of a prompt than the middle.
- Clearly separate instructions from the text they apply to - more on this in the next section.
- Be specific and descriptive about the task and the desired outcome - its format, length, style, language, etc.
- Avoid ambiguous descriptions and instructions.
- Favor instructions that say “what to do” instead of those that say “what not to do”.
- “Lead” the output in the right direction by writing the first word (or even begin the first sentence for the model).
- Use advanced techniques like Few-shot prompting and Chain-of-thought
- Test your prompts with different models to assess their robustness.
- Version and track the performance of your prompts.



Hugging Face

<https://huggingface.co/docs/transformers/main/tasks/prompting#best-practices-of-llm-prompting>


“Create three ideas for social media content”



- Not enough context (e.g., which social media platform)
- Not enough details about the content (e.g., topics, audience)

“Compose a thought-provoking LinkedIn post about the future of remote work. Please write in the first person, highlight the advantages of remote collaboration, mention industry trends, and write for young managers working in the tech sector.”

<https://www.planthat.com/good-vs-bad-ai-prompts/>



Mentimeter



Prompt Engineering

Here is another example of a bad prompt:

“Create a social media post about the lecture I attended at university today”


Please give an example of a better prompt.

Challenges and Risks of Generative AI



A hallucination is when a language model generates seemingly realistic responses that are untrue, nonsensical, or unfaithful to the provided source input.

- LLMs hallucinate because they operate mainly on a probabilistic level by putting one probable word in front of another without any deeper capacity for comprehension → **“Stochastic Parrots”**
- Like a parrot, LLMs excel at regurgitating learned content without knowing the meaning of their responses



Hannigan et al. 2024, Bender et al. 2021



Botshit is AI-generated content that is not grounded in truth (e.g., hallucinations) and is then uncritically used by a human for communication and decision-making tasks.

- Humans might use untrue material created by generative AI in an uncritical and thoughtless way
- This “botshit” makes it harder for people to know what is true and false in the world
- Problems arise when the outputs have important consequences and cannot easily be verified



Hannigan et al. 2024; McCarthy et al. 2024

New York lawyers sanctioned for using fake ChatGPT cases in legal brief

NEW YORK, June 22 (Reuters) - A U.S. judge on Thursday imposed sanctions on two New York lawyers who submitted a legal brief that included six fictitious case citations generated by an artificial intelligence chatbot, ChatGPT.

<https://www.reuters.com/legal/new-york-lawyers-sanctioned-using-fake-chatgpt-cases-legal-brief-2023-06-22/>

Sports Illustrated Published Articles by Fake, AI-Generated Writers

There was nothing in Drew Ortiz's [author biography](#) at *Sports Illustrated* to suggest that he was anything other than human.

"Drew has spent much of his life outdoors, and is excited to guide you through his never-ending list of the best products to keep you from falling to the perils of nature," it read. "Nowadays, there is rarely a weekend that goes by where Drew isn't out camping, hiking, or just back on his parents' farm."

The only problem? Outside of *Sports Illustrated*, Drew Ortiz doesn't seem to exist. He has no social media presence and no publishing history. And even more strangely, his profile photo on *Sports Illustrated* [is for sale](#) on a website that sells AI-generated headshots, where he's described as "neutral white young-adult male with short brown hair and blue eyes."

<https://futurism.com/sports-illustrated-ai-generated-writers>

- The data used to train AI generative AI models often reflect human biases:
 - Cultural biases
 - Gender biases
 - Racial biases
 - Socioeconomic biases
 - ...
- Like any AI, these models inherit those biases


Generative AI: UNESCO study reveals alarming evidence of regressive gender stereotypes

Ahead of the International Women's Day, a UNESCO study revealed worrying tendencies in Large Language models (LLM) to produce gender bias, as well as homophobia and racial stereotyping. Women were described as working in domestic roles far more often than men – four times as often by one model – and were frequently associated with words like “home”, “family” and “children”, while male names were linked to “business”, “executive”, “salary”, and “career”.

<https://www.unesco.org/en/articles/generative-ai-unesco-study-reveals-alarming-evidence-regressive-gender-stereotypes>

Getty Images lawsuit says Stability AI misused photos to train AI

(Reuters) - Stock photo provider Getty Images has sued artificial intelligence company Stability AI Inc, accusing it in a lawsuit made public on Monday of misusing more than 12 million Getty photos to train its Stable Diffusion AI image-generation system.

The [lawsuit](#) , filed in Delaware federal court, follows a separate Getty case against Stability in the United Kingdom and a related class-action [complaint](#) filed by artists in California against Stability and other companies in the fast-growing field of generative AI.

<https://www.reuters.com/legal/getty-images-lawsuit-says-stability-ai-misused-photos-train-ai-2023-02-06/>

The Times Sues OpenAI and Microsoft Over A.I. Use of Copyrighted Work

The New York Times sued OpenAI and Microsoft for copyright infringement on Wednesday, opening a new front in the increasingly intense legal battle over the unauthorized use of published work to train artificial intelligence technologies.

The Times is the first major American media organization to sue the companies, the creators of ChatGPT and other popular A.I. platforms, over copyright issues associated with its written works. [The lawsuit, filed in Federal District Court in Manhattan](#), contends that millions of articles published by The Times were used to train automated chatbots that now compete with the news outlet as a source of reliable information.

<https://www.nytimes.com/2023/12/27/business/media/new-york-times-open-ai-microsoft-lawsuit.html>



<https://www.youtube.com/watch?v=LDfbGk9dpWw>

*“If you have a hammer,
everything looks like a nail”*

Law of the Instrument or Maslow’s Hammer (1966)

Use-case family	Generative models' current usefulness	Example use cases
Prediction/forecasting	Low	Risk prediction, customer churn prediction, sales/demand forecasting
Decision intelligence	Low	Decision support, augmentation, automation
Segmentation/classification	Medium	Clustering, customer segmentation, object classification
Recommendation systems	Medium	Recommendation engine, personalized advice, next best action
Content generation	High	Text generation, image and video generation, synthetic data
Conversational user interfaces	High	Virtual assistant, chatbot, digital worker

Source: Gartner
© 2024 Gartner, Inc. and/or its affiliates. All rights reserved. 2937191

Gartner


<https://www.gartner.com/en/articles/when-not-to-use-generative-ai>

Key Takeaways From This Lecture

- The primary way in which generative AI differs from traditional AI is that it can generate new data instances (new text, images, videos, etc.)
- Foundation models are at the core of generative AI
- Foundation models can be adapted to a wide range of specific tasks, for example, via fine-tuning or prompting
- There are different prompting techniques (zero-shot, few-shot, chain-of-thought prompting) and a growing number of guidelines for effective prompt design
- Generative AI offers many benefits but also creates several challenges such as hallucinations, biases, and copyright issues



***Thank you for
your attention!***



- Appel, G., Neelbauer, J., & Schweidel, D. A. (2023). Generative AI has an intellectual property problem. Harvard Business Review. <https://hbr.org/2023/04/generative-ai-has-an-intellectual-property-problem>
- Bender, E. M., Gebru, T., McMillan-Major, A., & Shmitchell, S. (2021, March). On the dangers of stochastic parrots: Can language models be too big? 🦜. In Proceedings of the 2021 ACM conference on fairness, accountability, and transparency (pp. 610-623).
- Bommasani, R., Hudson, D. A., Adeli, E., Altman, R., Arora, S., von Arx, S., ... & Liang, P. (2021). On the opportunities and risks of foundation models. arXiv preprint arXiv:2108.07258.
- Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., ... & Amodei, D. (2020). Language models are few-shot learners. Advances in neural information processing systems, 33, 1877-1901.
- EU AI Act. (2024). <http://data.europa.eu/eli/reg/2024/1689/oj>
- Faklaris, Cori. (2023). An Introduction to Generative AI. <https://corifaklaris.com/files/genAIintro.pdf>
- Feuerriegel, S., Hartmann, J., Janiesch, C., & Zschech, P. (2024). Generative AI. Business & Information Systems Engineering. 66, 111–126 (2024). <https://doi.org/10.1007/s12599-023-00834-7>
- Gartner 2024. <https://www.gartner.com/en/newsroom/press-releases/2024-08-21-gartner-2024-hype-cycle-for-emerging-technologies-highlights-developer-productivity-total-experience-ai-and-security>
- Google. (2023). Introduction to Generative AI. <https://www.youtube.com/watch?v=G2fqAlgmoPo>
- Hannigan, T., McCarthy, I. P., & Spicer, A. (2023). Beware of botshit: How to manage the epistemic risks of generative chatbots. Business Horizons, 67(5), 471-486.
- Maslow, A. H. (1966). The Psychology of Science. Harper & Row.
- McCarthy, I. P., Hannigan, T., & Spicer, A. (2023). The Risks of Botshit. Harvard Business Review. <https://hbr.org/2024/07/the-risks-of-botshit>
- Stollnitz, Bea. (2023). <https://bea.stollnitz.com/blog/how-gpt-works/>
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... & Polosukhin, I. (2017). Attention is all you need. Advances in neural information processing systems, 30.
- Zamfirescu-Pereira, J. D., Wong, R. Y., Hartmann, B., & Yang, Q. (2023, April). Why Johnny can't prompt: how non-AI experts try (and fail) to design LLM prompts. In Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems (pp. 1-21).